

Manuel Romero

Python ML with sklearn

Code bits taken from Book and Github

https://github.com/kjmazidi/Machine_Learning_2nd_edition/tree/master/Part_7_Neural_Networks

Reading the CSV file Auto.csv Part 1

Selecting Auto.csv from local drive

Double-click (or enter) to edit

```
import pandas as pd
from google.colab import files
uploaded = files.upload()
import io
df = pd.read_csv(io.BytesIO(uploaded['Auto.csv']))
# Dataset is now stored in a Pandas Dataframe
```

Choose Files Auto.csv

- **Auto.csv**(text/csv) - 17859 bytes, last modified: 4/10/2023 - 100% done
- Saving Auto.csv to Auto (4).csv

file is read and now we will output the first rows and dimensions of the table

`df.head()`

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin	name
0	18.0	8	307.0	130	3504	12.0	70.0	1	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70.0	1	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70.0	1	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70.0	1	amc rebel sst
4	17.0	8	302.0	140	3449	NaN	70.0	1	ford torino

```
print("The dimension of the Auto.csv data is : " , df.shape)
```

```
The dimension of the Auto.csv data is : (392, 9)
```

Data exploration with code Part 2

MPG:

Range: 9 - 46.6

AVG: 23.445918

```
df.mpg.describe()
```

```
count    392.000000
mean     23.445918
std       7.805007
min       9.000000
25%      17.000000
50%      22.750000
75%      29.000000
max      46.600000
Name: mpg, dtype: float64
```

Weight:

Range: 1613 - 5140

AVG: 2977.584184

```
df.weight.describe()
```

```
count    392.000000
mean     2977.584184
std      849.402560
min      1613.000000
25%     2225.250000
50%     2803.500000
75%     3614.750000
max      5140.000000
Name: weight, dtype: float64
```

Year :

Range: 70 - 82

AVG: 76.010256

```
df.year.describe()
```

```
count    390.000000
mean      76.010256
```

```
std      3.668093
min      70.000000
25%      73.000000
50%      76.000000
75%      79.000000
max      82.000000
Name: year, dtype: float64
```

Explore Data types part 3

data types of columns

```
df.dtypes
```

```
mpg      float64
cylinders    int64
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year         float64
origin        int64
name         object
dtype: object
```

change the cylinders column to categorical (use cat.codes)

```
df.cylinders = df.cylinders.astype("category").cat.codes
```

change the origin column to categorical (don't use cat.codes)

```
df.origin = df.origin.astype("category")
```

verifying the changes with the dtypes attribute

```
df.dtypes
```

```
mpg      float64
cylinders    int8
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year         float64
origin        category
name         object
dtype: object
```

Deal with NAs part 4

```
df = df.dropna()
```

Dropped NA's

New dimensions Dimension 389 , 9

```
df.shape

(389, 9)
```

Modify columns part 5

Make a new column, mpg_high, and make it categorical

adding new blank column and then filling it in

```
##new column mpg_high set with initial value of 0
df['mpg_high'] = 0
```

Checking to see if column added

```
df.dtypes

mpg                float64
cylinders           int8
displacement       float64
horsepower         int64
weight             int64
acceleration       float64
year              float64
origin             category
name              object
mpg_high           int64
dtype: object
```

```
##getting average of mpg
avg = df["mpg"].mean()
print(avg)
##filling in the column with 1 if mpg > mean mpg else is covered since we made the new column with 0s originally
df['mpg_high'][df.mpg > avg] = 1
```

```
23.490488431876607
<ipython-input-114-596124658d86>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['mpg_high'][df.mpg > avg] = 1
```

```
## making mpg high a category
df.mpg_high = df.mpg_high.astype("category")
```

deleting mpg and name column and printing df.head

```
df = df.drop(columns=["mpg", "name"])
print(df.head())
```

	cylinders	displacement	horsepower	weight	acceleration	year	origin	\
0	4	307.0	130	3504	12.0	70.0	1	
1	4	350.0	165	3693	11.5	70.0	1	
2	4	318.0	150	3436	11.0	70.0	1	
3	4	304.0	150	3433	12.0	70.0	1	
6	4	454.0	220	4354	9.0	70.0	1	

	mpg_high
0	0
1	0
2	0
3	0
6	0

Data exploration with graphs part 6

catplot

```
import seaborn as sb
from sklearn import datasets
sb.catplot(x="mpg_high", kind="count", data=df)
```


```
<seaborn.axisgrid.FacetGrid at 0x7f983e9708b0>
```



The catplot shows us that there is more vehicles with a high mpg than those with not a high mpg.

```
sb.relplot(x='horsepower', y='weight', data=df, hue=df.mpg_high, style=df.mpg_high)
```


```
<seaborn.axisgrid.FacetGrid at 0x7f98463120d0>
```



We can see that the lower the weight the lower the horsepower, this makes sense since it takes less power to move less weight also less power typically means smaller engine. Furthermore we see that most lower weight and lower powered cars are high_mpg cars this also makes sense it takes less gas to move the weight and smaller engines use less gas.

```
sb.boxplot(x = 'mpg_high', y = 'weight', data = df)
```

```
<Axes: xlabel='mpg_high', ylabel='weight'>
```



once again we see that the high mpg car are lower in weight, the mean for high mpg is much lower than that of not high mpg. We also see that there are possible outliers on the data for high mpg vehicles. Some heavy cars are still getting high mpg.

Train/test split part 7

```
df.shape
```

```
(389, 8)
```

```
df.head()
```

	cylinders	displacement	horsepower	weight	acceleration	year	origin	mpg_high
0	4	307.0	130	3504	12.0	70.0	1	0
1	4	350.0	165	3693	11.5	70.0	1	0
2	4	318.0	150	3436	11.0	70.0	1	0
3	4	304.0	150	3433	12.0	70.0	1	0
6	4	454.0	220	4354	9.0	70.0	1	0



```
# train test split
from sklearn.model_selection import train_test_split
X = df.iloc[:, 0:6]
##7 is the mpg_high column
y = df.iloc[:, 7]
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=1234)
print('train size:', X_train.shape)
print('test size:', X_test.shape)

train size: (311, 6)
test size: (78, 6)
```

Logistic Regression pt 8

train a logistic regression model using solver lbfgs

test and evaluate

print metrics using the classification report

```
from sklearn.linear_model import LogisticRegression
clf = LogisticRegression()
clf.fit(X_train, y_train)
clf.score(X_train, y_train)
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
0.8102893890675241
```

output:

0.8102893890675241

```
# make predictions
pred = clf.predict(X_test)
from sklearn.metrics import classification_report
print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	0.85	0.74	0.79	39
1	0.77	0.87	0.82	39
accuracy			0.81	78

macro avg	0.81	0.81	0.81	78
weighted avg	0.81	0.81	0.81	78

Decision Tree part 9

```
from sklearn import tree
clf = DecisionTreeClassifier()
clf.fit(X_train, y_train)
pred = clf.predict(X_test)
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
print('accuracy score: ', accuracy_score(y_test, pred))
print('precision score: ', precision_score(y_test, pred))
print('recall score: ', recall_score(y_test, pred))
print('f1 score: ', f1_score(y_test, pred))
```

```
accuracy score: 0.7307692307692307
precision score: 0.68
recall score: 0.8717948717948718
f1 score: 0.7640449438202247
```

```
from sklearn.metrics import confusion_matrix
confusion_matrix(y_test, pred)
from sklearn.metrics import classification_report
print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	0.82	0.59	0.69	39
1	0.68	0.87	0.76	39
accuracy			0.73	78
macro avg	0.75	0.73	0.73	78
weighted avg	0.75	0.73	0.73	78

We see that logistic regression had better accuracy than that of decision trees.

```
tree.plot_tree(clf)
```

```
[Text(0.46629213483146065, 0.9583333333333334, 'x[1] <= 153.0\ngini = 0.447\nsamples = 311\nvalue = [105, 206]'),  
Text(0.2303370786516854, 0.875, 'x[3] <= 2822.5\ngini = 0.11\nsamples = 171\nvalue = [10, 161]'),  
Text(0.16853932584269662, 0.7916666666666666, 'x[5] <= 73.5\ngini = 0.064\nsamples = 152\nvalue = [5, 147]'),  
Text(0.14606741573033707, 0.7083333333333334, 'x[1] <= 131.0\ngini = 0.239\nsamples = 36\nvalue = [5, 31]'),  
Text(0.0898876404494382, 0.625, 'x[0] <= 0.5\ngini = 0.117\nsamples = 32\nvalue = [2, 30]'),  
Text(0.0449438202247191, 0.5416666666666666, 'x[5] <= 72.5\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),  
Text(0.02247191011235955, 0.4583333333333333, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),  
Text(0.06741573033707865, 0.4583333333333333, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),  
Text(0.1348314606741573, 0.5416666666666666, 'x[3] <= 2377.0\ngini = 0.064\nsamples = 30\nvalue = [1, 29]'),  
Text(0.11235955056179775, 0.4583333333333333, 'gini = 0.0\nsamples = 25\nvalue = [0, 25]'),  
Text(0.15730337078651685, 0.4583333333333333, 'x[3] <= 2387.0\ngini = 0.32\nsamples = 5\nvalue = [1, 4]'),  
Text(0.1348314606741573, 0.375, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),  
Text(0.1797752808988764, 0.375, 'gini = 0.0\nsamples = 4\nvalue = [0, 4]'),  
Text(0.20224719101123595, 0.625, 'x[3] <= 2332.5\ngini = 0.375\nsamples = 4\nvalue = [3, 1]'),  
Text(0.1797752808988764, 0.5416666666666666, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),  
Text(0.2247191011235955, 0.5416666666666666, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),  
Text(0.19101123595505617, 0.7083333333333334, 'gini = 0.0\nsamples = 116\nvalue = [0, 116]'),  
Text(0.29213483146067415, 0.7916666666666666, 'x[5] <= 79.5\ngini = 0.388\nsamples = 19\nvalue = [5, 14]'),  
Text(0.2696629213483146, 0.7083333333333334, 'x[2] <= 87.5\ngini = 0.496\nsamples = 11\nvalue = [5, 6]'),  
Text(0.24719101123595505, 0.625, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),  
Text(0.29213483146067415, 0.625, 'x[5] <= 75.5\ngini = 0.469\nsamples = 8\nvalue = [5, 3]'),  
Text(0.2696629213483146, 0.5416666666666666, 'x[3] <= 2900.5\ngini = 0.375\nsamples = 4\nvalue = [1, 3]'),  
Text(0.24719101123595505, 0.4583333333333333, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),  
Text(0.29213483146067415, 0.4583333333333333, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),  
Text(0.3146067415730337, 0.5416666666666666, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),  
Text(0.3146067415730337, 0.7083333333333334, 'gini = 0.0\nsamples = 8\nvalue = [0, 8]'),  
Text(0.702247191011236, 0.875, 'x[5] <= 75.5\ngini = 0.436\nsamples = 140\nvalue = [95, 45]'),  
Text(0.5168539325842697, 0.7916666666666666, 'x[4] <= 14.75\ngini = 0.289\nsamples = 80\nvalue = [66, 14]'),  
Text(0.4044943820224719, 0.7083333333333334, 'x[2] <= 151.5\ngini = 0.113\nsamples = 50\nvalue = [47, 3]'),  
Text(0.38202247191011235, 0.625, 'x[4] <= 10.25\ngini = 0.198\nsamples = 27\nvalue = [24, 3]'),  
Text(0.3595505617977528, 0.5416666666666666, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),  
Text(0.4044943820224719, 0.5416666666666666, 'x[3] <= 4062.0\ngini = 0.142\nsamples = 26\nvalue = [24, 2]'),  
Text(0.38202247191011235, 0.4583333333333333, 'gini = 0.0\nsamples = 15\nvalue = [15, 0]'),  
Text(0.42696629213483145, 0.4583333333333333, 'x[3] <= 4089.0\ngini = 0.298\nsamples = 11\nvalue = [9, 2]'),  
Text(0.4044943820224719, 0.375, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),  
Text(0.449438202247191, 0.375, 'x[3] <= 4448.5\ngini = 0.18\nsamples = 10\nvalue = [9, 1]'),  
Text(0.42696629213483145, 0.2916666666666667, 'gini = 0.0\nsamples = 6\nvalue = [6, 0]'),  
Text(0.47191011235955055, 0.2916666666666667, 'x[3] <= 4460.5\ngini = 0.375\nsamples = 4\nvalue = [3, 1]'),  
Text(0.449438202247191, 0.2083333333333334, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),  
Text(0.4943820224719101, 0.2083333333333334, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),  
Text(0.42696629213483145, 0.625, 'gini = 0.0\nsamples = 23\nvalue = [23, 0]'),  
Text(0.6292134831460674, 0.7083333333333334, 'x[2] <= 125.0\ngini = 0.464\nsamples = 30\nvalue = [19, 11]'),  
Text(0.6067415730337079, 0.625, 'x[3] <= 3111.5\ngini = 0.393\nsamples = 26\nvalue = [19, 7]'),  
Text(0.5393258426966292, 0.5416666666666666, 'x[2] <= 92.5\ngini = 0.444\nsamples = 9\nvalue = [3, 6]'),  
Text(0.5168539325842697, 0.4583333333333333, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),  
Text(0.5617977528089888, 0.4583333333333333, 'x[4] <= 15.25\ngini = 0.245\nsamples = 7\nvalue = [1, 6]'),  
Text(0.5393258426966292, 0.375, 'x[2] <= 105.0\ngini = 0.5\nsamples = 2\nvalue = [1, 1]'),  
Text(0.5168539325842697, 0.2916666666666667, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),  
Text(0.5617977528089888, 0.2916666666666667, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),  
Text(0.5842696629213483, 0.375, 'gini = 0.0\nsamples = 5\nvalue = [0, 5]'),  
Text(0.6741573033707865, 0.5416666666666666, 'x[3] <= 3384.0\ngini = 0.111\nsamples = 17\nvalue = [16, 1]'),  
Text(0.651685393258427, 0.4583333333333333, 'x[3] <= 3332.5\ngini = 0.219\nsamples = 8\nvalue = [7, 1]'),  
Text(0.6292134831460674, 0.375, 'gini = 0.0\nsamples = 7\nvalue = [7, 0]'),  
Text(0.6741573033707865, 0.375, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),  
Text(0.6966292134831461, 0.4583333333333333, 'gini = 0.0\nsamples = 9\nvalue = [9, 0]'),  
Text(0.651685393258427, 0.625, 'gini = 0.0\nsamples = 4\nvalue = [0, 4]'),  
Text(0.8876404494382022, 0.7916666666666666, 'x[5] <= 79.5\ngini = 0.499\nsamples = 60\nvalue = [29, 31]'),
```

```
Text(0.8426966292134831, 0.7083333333333334, 'x[4] <= 18.95\ngini = 0.493\nsamples = 50\nvalue = [28, 22]'),
```

Here is the decision Tree plotted

```
Text(0.7415730337078652, 0.4583333333333333, 'x[3] <= 3507.5\ngini = 0.498\nsamples = 15\nvalue = [8, 7] ),
```

Neural Network Part 10

```
Text(0.7415730337078652, 0.4583333333333333, 'gini = 0.0\nsamples = 5\nvalue = [5, 0] ),
```

Neural Network

```
Text(0.7065168538325842, 0.125, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
```

```
from sklearn import preprocessing
```

```
scaler = preprocessing.StandardScaler().fit(X_train)
```

```
X_train_scaled = scaler.transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
Text(0.8539325842696629, 0.375, 'x[1] <= 237.5\ngini = 0.375\nsamples = 4\nvalue = [3, 1]'),
```

```
# train
```

```
from sklearn.neural_network import MLPClassifier
```

```
clf = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(5, 2), max_iter=500, random_state=1234)
```

```
clf.fit(X_train_scaled, y_train)
```

```
▼
MLPClassifier
MLPClassifier(hidden_layer_sizes=(5, 2), max_iter=500, random_state=1234,
               solver='lbfgs')
```

```
Text(0.9775280898876404, 0.5416666666666666, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
```

```
pred = clf.predict(X_test_scaled)
```

```
# output results
```

```
print('accuracy = ', accuracy_score(y_test, pred))
```

```
confusion_matrix(y_test, pred)
```

```
from sklearn.metrics import classification_report
```

```
print(classification_report(y_test, pred))
```

```
accuracy = 0.8333333333333334
precision    recall  f1-score   support

0           0.80     0.90     0.84         39
1           0.88     0.77     0.82         39

accuracy          0.83         78
macro avg         0.84         0.83         0.83         78
weighted avg      0.84         0.83         0.83         78
```



different settings

```

clf = MLPClassifier(solver='sgd', hidden_layer_sizes=(4,6), max_iter=1500, random_state=1234)
clf.fit(X_train_scaled, y_train)
# make predictions
pred = clf.predict(X_test_scaled)

print('accuracy = ', accuracy_score(y_test, pred))

# confusion matrix
confusion_matrix(y_test, pred)
print(classification_report(y_test, pred))

```

```

accuracy = 0.8333333333333334
          precision    recall  f1-score   support

     0         0.86      0.79      0.83         39
     1         0.81      0.87      0.84         39

 accuracy                   0.83         78
  macro avg              0.84         0.83         0.83         78
 weighted avg            0.84         0.83         0.83         78

```

using the neural network with `clf = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(5, 2), max_iter=500, random_state=1234)` provides a much better accuracy than logistic regression and about the same for second model of neural networks. However the first model works better for precision on high mpg vs the second one being better for not high mpg classification. The reason for this is the change in hidden layers.

My preference was python it feels much better to use and much faster. One issue was that my computer isn't updating my pandas and therefore is not working great but with google colab that was solved and everything went smoothly after that. On the other hand R sometimes installs packages successfully and others it won't not sure why, but probably user error or machine error that I need to look into.

My preference was python it feels much better to use and much faster. One issue was that my computer isn't updating my pandas and therefore is not working great but with google colab that was solved and everything went smoothly after that. On the other hand R sometimes installs packages successfully and others it won't not sure why, but probably user error or machine error that I need to look into.

✓ 1s completed at 4:17 PM

